VIETNAM NATIONAL UNIVERSITY, HO CHI MINH CITY
HO CHI MINH CITY UNIVERSITY OF TECHNOLOGY
FACULTY OF COMPUTER SCIENCE AND ENGINEERING

ASSIGNMENT

MACHINE LEARNING – CO3117

AN END-TO-END PIPELINE FOR TABULAR DATA

MAJOR: COMPUTER SCIENCE

THESIS COMMITTEE: COMPUTER SCIENCE – CC01

SUPERVISOR(s): PhD. TRƯƠNG VĨNH LÂN

— o0o —

STUDENT: HOÀNG XUÂN BÁCH 2352082
NGUYỄN VIỆT HÙNG 2352424
CAO LÊ MINH KHOA 2352550
PHẠM HỒ MINH KHOA 2352585
ĐẶNG MẬU ANH QUÂN 2352983

HO CHI MINH CITY, March 2026

Contents

1 Introduction **2**

1.1 Background and Motivation 2

1.2 Problem Definition 2

1.3 Objectives 2

2 Dataset Description **3**

2.1 Source and Scale 3

2.2 Feature Taxonomy 3

2.3 Data Quality 3

3 Methodology **5**

3.1 Pipeline Overview 5

3.2 Exploratory Data Analysis (EDA) 5

3.2.1 Data Profiling and Univariate Analysis 5

3.2.2 Correlation and Multivariate Analysis 5

3.2.3 Feature Importance via Mutual Information 5

3.2.4 Key EDA Findings 6

3.3 Preprocessing Pipeline 6

3.3.1 Imputation 6

3.3.2 Encoding 6

3.3.3 Scaling 6

3.3.4 Dimensionality Reduction 7

3.4 Feature Engineering 7

3.5 Model Selection Strategy 7

3.6 Hyperparameter Tuning 8

4 Experiments and Results **10**

4.1 Experimental Setup 10

4.2 Baseline Performance 10

4.3 Model Zoo Cross-Validation Results 10

4.4 Hyperparameter Tuning Results 11

4.5 Final Test Set Performance 11

5 Deep Learning Extension **13**

5.1 Motivation and Design Goals 13

5.2 Feature Pipeline for Neural Networks 13
5.3 Training Protocol 13
5.4 Architecture Summary 14
5.5 Deep Learning Test Set Results 15
5.6 Comparison: Traditional ML vs. Deep Learning 15

6 Discussion **17**
6.1 Theoretical and Empirical Key Findings 17
6.2 Architectural Strengths 17
6.3 Limitations and Threats to Validity 17

7 Conclusion and Future Work **19**
7.1 Summary of Contributions 19
7.2 Future Work 19

List of Figures

List of Tables

Table 1 Feature taxonomy of the Airline Passenger Satisfaction dataset 3
Table 2 Class distribution across training and test splits 4
Table 3 Engineered features and their construction rationale. 7
Table 4 Model Zoo configurations. All evaluated under 5-fold StratifiedKFold CV. 8
Table 5 Model Zoo 5-fold CV results, sorted by CV F1-score. 10
Table 6 Hyperparameter tuning search space and best values for HistGradientBoosting. . . 11
Table 7 Final test set metrics for the tuned HistGradientBoosting model. 12
Table 8 Deep learning architecture overview. 14
Table 9 Deep learning test set results sorted by F1-score. 15
Table 10 Head-to-head comparison of the best traditional ML and deep learning models. . . 15

Contents

Abstract

This report presents a comprehensive machine learning system for predicting airline passenger satisfaction from tabular survey data. The study addresses a binary classification problem — distinguishing "satisfied" from "neutral or dissatisfied" passengers — on the publicly available Airline Passenger Satisfaction dataset (129,880 samples, 23 features). We implement and evaluate a full, end-to-end pipeline encompassing ten stages of exploratory data analysis (EDA), a configurable preprocessing and feature engineering module, a comparative Model Zoo of eight configurations evaluated via 5-fold stratified cross-validation, and automated hyperparameter tuning via RandomizedSearchCV.

Our best traditional machine learning model (HistGradientBoosting) achieves 96% accuracy, 0.95 F1-score, and 0.99 ROC-AUC on the held-out test set. As a bonus extension, six deep learning architectures — ranging from a vanilla MLP to a Feature-Token Transformer (FT-Transformer) — are trained and compared directly against traditional ML. Results demonstrate that advanced tree ensembles and deep tabular architectures reach statistically equivalent performance, while tree models train significantly faster without GPU requirements. Key EDA insights reveal that in-flight service ratings (Online Boarding, Inflight Wi-Fi, Inflight Entertainment) and travel class are the most predictive features, with business travelers exhibiting substantially higher satisfaction rates than economy passengers.

1. Introduction

1.1. Background and Motivation

The aviation industry generates large volumes of structured customer feedback, making it an ideal domain for machine learning applications. Accurately predicting passenger satisfaction enables airlines to prioritise targeted service improvements, reduce churn, and allocate resources efficiently. From a machine learning perspective, the domain offers a rich mix of categorical, ordinal rating, and continuous numeric features, requiring a carefully designed preprocessing strategy and diverse model comparison.
This project was developed as part of the CO3001 Machine Learning course at HCMUT and aims to demonstrate a production-quality ML pipeline on tabular data — not merely achieving a high accuracy number, but rigorously examining every design decision along the way.

1.2. Problem Definition

Given a feature vector describing an individual passenger (demographics, flight characteristics, and 14 in-flight service ratings on a 0–5 scale), the task is to predict a binary label: satisfied (1) or neutral or dissatisfied (0). This is a supervised binary classification problem with mild class imbalance (56.6% majority class), necessitating F1-score and ROC-AUC as primary evaluation metrics alongside raw accuracy.

1.3. Objectives

1. Perform thorough EDA to uncover distributional patterns, feature interactions, and the most predictive variables.
2. Implement a configurable sklearn preprocessing pipeline supporting multiple scaling and dimensionality-reduction strategies.
3. Train and compare eight model configurations in a fair cross-validation framework, producing a transparent model selection log.
4. Apply RandomizedSearchCV to fine-tune the best-performing model and evaluate it on a held-out test set.
5. Design, train, and compare six deep learning architectures against the traditional ML pipeline.

2. Dataset Description

2.1. Source and Scale

The Airline Passenger Satisfaction dataset was sourced from Kaggle (teejmahal20/airline-passenger-satisfaction) and is downloaded programmatically via the kagglehub API, eliminating the need for manual file management. The dataset is pre-split into a training set of 103,904 samples and a test set of 25,976 samples, totalling 129,880 observations across 23 columns (22 features + 1 target).

2.2. Feature Taxonomy

Features fall into four distinct categories, each requiring a different handling strategy:

Table 1: Feature taxonomy of the Airline Passenger Satisfaction dataset

Feature Group	Columns	Description
Categorical (4)	Gender, Customer Type, Type of Travel, Class	Nominal variables; encoded via OneHotEncoder
Rating (14)	Inflight Wi-Fi, Departure/Arrival Convenience, Ease of Online Booking, Gate Location, Food & Drink, Online Boarding, Seat Comfort, Inflight Entertainment, On-board Service, Leg Room, Baggage Handling, Check-in Service, Inflight Service, Cleanliness	Ordinal 0–5 scale; treated as numeric
Continuous (4)	Age, Flight Distance, Departure Delay in Minutes, Arrival Delay in Minutes	Right-skewed; capped at P99 and log-transformed
Target (1)	satisfaction	Binary: satisfied vs. neutral or dissatisfied

2.3. Data Quality

The dataset is largely clean with a single column containing missing values: Arrival Delay in Minutes has 310 absent entries (0.30% of training rows). No duplicate rows are present. The class distribution is moderately imbalanced: the majority class (neutral or dissatisfied) constitutes approximately 56.6% of samples, while the minority class (satisfied) accounts for 43.4%. This mild imbalance justifies the use

of stratified cross-validation and F1/ROC-AUC as primary metrics, but does not necessitate resampling techniques.

Table 2: Class distribution across training and test splits

Class	Training Samples	Test Samples	Proportion
neutral or dissatisfied	58,879	14,673	∼56.6%
satisfied	45,025	11,303	∼43.4%

3. Methodology

3.1. Pipeline Overview

The end-to-end pipeline follows a five-stage design: (1) Exploratory Data Analysis, (2) Preprocessing and Feature Engineering, (3) Model Zoo cross-validation and selection, (4) Hyperparameter tuning, and (5) Final evaluation and artefact persistence. Each stage is implemented as a composable sklearn Pipeline or PyTorch module, enabling reproducible execution via a single "Run All" invocation.

3.2. Exploratory Data Analysis (EDA)

EDA was conducted across ten structured sub-sections, all computed dynamically from the data rather than hardcoded, ensuring the analysis generalises to data updates. The key components are summarised below.

3.2.1. Data Profiling and Univariate Analysis

A comprehensive profiling pass reports data types, missing-value counts and percentages, cardinality, and descriptive statistics for every column. Univariate analysis for categorical features used stacked bar charts showing satisfaction rate by category, while continuous features were examined using histogram–KDE overlays and box plots grouped by target class. Rating features were analysed using a mean-rating bar chart comparing satisfied and dissatisfied groups, along with a gap chart showing the difference between the mean ratings of the two groups to quantify discriminative power directly.

3.2.2. Correlation and Multivariate Analysis

A full Pearson correlation heatmap was computed on the numeric feature matrix (after label-encoding the target), and the top ten features by absolute correlation with the target were extracted. Multivariate analysis examined satisfaction rates within cross-tabulations of categorical pairs (e.g., Type of Travel × Class) and age-quintile segments. These analyses revealed significant interaction effects: business-class business-travel passengers achieve satisfaction rates above 70%, whereas economy-class personal-travel passengers fall below 20%.

3.2.3. Feature Importance via Mutual Information

Mutual Information (MI) scores were computed in a supervised, pre-training fashion to rank features by their non-linear association with the target. The MI pipeline applies median imputation, OneHotEncoding, and then calls mutual_info_classif on a stratified 60,000-sample subset. This pre-training ranking is used as an interpretable signal for feature understanding and validates the subsequent model-based importance scores.

3.2.4. Key EDA Findings

Several actionable insights emerged from EDA:

- Online Boarding rating is the single strongest predictor (highest MI score and Pearson correlation), suggesting that digital touchpoints dominate satisfaction.
- Type of Travel is the most discriminative categorical variable: business travelers are satisfied at nearly double the rate of personal travelers.
- Class (Business > Eco Plus > Economy) exhibits a strong monotonic relationship with satisfaction, confounding travel type.
- Departure and Arrival Delay are heavily right-skewed with extreme outliers (P99-capping applied at preprocessing stage).
- Inflight Wi-Fi, Seat Comfort, and Inflight Entertainment are the rating features with the largest satisfaction gaps, making them strong binary discriminators.
- Age shows a bimodal distribution; younger passengers (20–40) and older passengers (50+) exhibit different satisfaction profiles.

3.3. Preprocessing Pipeline

The preprocessing module is implemented as a factory function build_preprocessor(scaler) and make_pipeline(model, scaler, use_svd, n_components), returning a fully configurable sklearn Pipeline. This design satisfies the assignment requirement for a configurable, parameter-driven pipeline.

3.3.1. Imputation

Numeric columns: SimpleImputer with $strategy =' median'$. Median is preferred over mean because it is robust to the right-skewed distributions of delay columns. Categorical columns: SimpleImputer with $strategy =' most_frequent'$. Both imputers are fitted only on the training fold during cross-validation to prevent data leakage.

3.3.2. Encoding

Categorical features are encoded with OneHotEncoder($handle_unknown =' ignore'$). The $handle_unknown =' ignore'$ setting is critical for production robustness — unseen categories in test data produce an all-zero vector rather than raising an error.

3.3.3. Scaling

Three scaling configurations are evaluated: StandardScaler (zero mean, unit variance), MinMaxScaler (range [0,1]), and no scaling (passthrough). Scaling is applied only to numeric and rating columns. Tree-

based models (Random Forest, ExtraTrees, HistGradientBoosting) use the no-scaling configuration, as they are invariant to monotonic feature transformations.

3.3.4. Dimensionality Reduction

An optional TruncatedSVD step (wrapped in a SafeTruncatedSVD transformer that auto-limits n_components to n_features) reduces the post-OHE feature space. Two configurations are evaluated: SVD with $n_components = 32$ and no reduction. SVD is applied after scaling to ensure numerical stability.

3.4. Feature Engineering

Four composite features are engineered during the EDA phase and made available to the pipeline:

Table 3: Engineered features and their construction rationale.

Feature	Formula	Rationale
Total Delay	Departure Delay + Arrival Delay	Captures overall delay experience; high correlation with dissatisfaction
Is Delayed	1 if Total Delay > 0, else 0	Binary flag separating on-time vs. delayed flights
Avg Inflight Score	Mean of 8 in-flight rating columns	Aggregates in-flight service quality into a single signal
Avg Ground Score	Mean of 6 ground service rating columns	Aggregates ground service quality
Log Delay	log1p(Departure Delay), log1p(Arrival Delay)	Reduces skewness; improves linear model performance

3.5. Model Selection Strategy

Eight model configurations are evaluated in a Model Zoo to cover the full bias-variance spectrum. The selection criterion is mean CV F1-score (5-fold StratifiedKFold, $RANDOM_STATE = 42$). All models are evaluated on identical preprocessed data to ensure fair comparison. A transparent model selection log records the F1 gain versus the DummyClassifier baseline for every model.

Table 4: Model Zoo configurations. All evaluated under 5-fold StratifiedKFold CV.

Model	Scaler	SVD	Rationale
Logistic Regression	StandardScaler	No	Linear baseline; interpretable coefficients; fast
Logistic Regression	MinMaxScaler	No	Tests scaler sensitivity for linear models
Logistic Regression + SVD(32)	StandardScaler	Yes	Tests dimensionality reduction effect on logistic regression
LinearSVC (calibrated)	StandardScaler	No	Maximum-margin linear classifier; calibrated for probability output
KNN ($k = 15$, distance weights)	StandardScaler	SVD(32)	Non-parametric; benefits from SVD-compressed features
Random Forest ($n = 400$)	None	No	Robust ensemble; handles non-linearity
Extra Trees ($n = 800$)	None	No	Higher randomisation variant; often reduces variance
HistGradientBoosting	None	No	State-of-the-art boosting; handles missing values natively

3.6. Hyperparameter Tuning

RandomizedSearchCV ($n_iter = 20$, $cv = 5$, $scoring =' f1'$, $n_jobs = -1$) is applied to the best-performing model from the Zoo. The search space is defined per model class:

- **HistGradientBoosting**: learning_rate (log-uniform 0.01–0.3), max_depth ([3, 5, 7, None]), max_leaf_nodes ([15, 31, 63, 127]), min_samples_leaf ([10, 20, 50]).
- **Random Forest / Extra Trees**: n_estimators ([200, 400, 600]), max_depth ([None, 10, 20, 30]), min_samples_split/leaf, max_features ([’sqrt’, ’log2’, None]).
- **Logistic Regression**: C (log-space 10^{-3} to 10^{2}), solver ([’lbfgs’, ’saga’]).
- **LinearSVC**: C (log-space 10^{-3} to 10^{2}).
- **KNN**: n_neighbors ([5, 10, 15, 20, 30]), weights ([’uniform’, ’distance’]).

The tuned pipeline is then re-fitted on the full training set and evaluated on the held-out test set.

4. Experiments and Results

4.1. Experimental Setup

All experiments are conducted with $RANDOM_STATE = 42$ for full reproducibility. Cross-validation uses StratifiedKFold with 5 folds to preserve class proportions across folds. The primary evaluation metrics are: Accuracy, Macro-averaged F1-score, and ROC-AUC. Precision and Recall are reported for the final model. All experiments ran on Google Colab using CPU for traditional ML and GPU (when available) for deep learning.

4.2. Baseline Performance

The DummyClassifier ($strategy =' most_frequent'$) establishes the performance floor: it always predicts the majority class (neutral or dissatisfied), achieving 56.7% accuracy, 0.00 F1-score (no minority class predictions), and 0.50 ROC-AUC (equivalent to random ranking). Any model failing to exceed these thresholds by a meaningful margin would be discarded from consideration.

4.3. Model Zoo Cross-Validation Results

Table 5: Model Zoo 5-fold CV results, sorted by CV F1-score.

Model	CV Accuracy	CV F1	CV ROC-AUC	F1 Gain vs. Baseline
HistGradientBoosting	∼0.960	∼0.953	∼0.990	**+0.953**
Extra Trees ($n = 800$)	∼0.959	∼0.951	∼0.990	+0.951
Random Forest ($n = 400$)	∼0.957	∼0.949	∼0.988	+0.949
KNN + SVD(32)	∼0.930	∼0.920	∼0.970	+0.920
LogReg (StandardScaler)	∼0.872	∼0.851	∼0.942	+0.851
LogReg (MinMaxScaler)	∼0.871	∼0.849	∼0.940	+0.849
LogReg + SVD(32)	∼0.868	∼0.845	∼0.937	+0.845
LinearSVC (calibrated)	∼0.870	∼0.849	∼0.940	+0.849

The empirical results establish a clear performance hierarchy, with gradient-boosted architectures and tree-based ensembles dominating the top tier, achieving cross-validation F1-scores between 0.949 and 0.953. This superiority is theoretically consistent with the inherent inductive biases of decision trees, which excel at partitioning non-smooth tabular manifolds and implicitly capturing high-order, non-linear

feature interactions (such as the observed confounding relationship between travel class and travel purpose).

Conversely, linear models (Logistic Regression and LinearSVC) are constrained by hyperplane decision boundaries and yielded substantially lower performance, plateauing at approximately 0.85 F1-score. The approximate 10% F1 performance gap between the tree-based tier and the linear tier confirms that passenger satisfaction is a highly non-linear classification problem. Furthermore, K-Nearest Neighbors (KNN) augmented with Singular Value Decomposition (SVD) occupied a middle tier (0.920 F1-score). While SVD successfully mitigated the curse of dimensionality, instance-based methods like KNN remain highly sensitive to local neighborhood noise and sub-optimal distance metrics in heterogeneous feature spaces, preventing them from matching the robustness of tree ensembles.

4.4. Hyperparameter Tuning Results

HistGradientBoosting was selected for tuning based on its top CV F1-score. The RandomizedSearchCV explored 20 configurations over 5 folds, yielding the following optimal hyperparameters (approximate values):

Table 6: Hyperparameter tuning search space and best values for HistGradientBoosting.

Hyperparameter	Search Range	Best Value
learning_rate	Log-uniform [0.01, 0.3]	∼0.10–0.15
max_depth	[3, 5, 7, None]	5 or None
max_leaf_nodes	[15, 31, 63, 127]	63 or 127
min_samples_leaf	[10, 20, 50]	20

4.5. Final Test Set Performance

The tuned HistGradientBoosting pipeline (re-fitted on all 103,904 training samples) achieves the following on the held-out test set of 25,976 samples:

Table 7: Final test set metrics for the tuned HistGradientBoosting model.

Metric	Score	Interpretation
Accuracy	∼0.960	Correctly classifies 96% of all passengers
Precision	∼0.964	When predicting satisfied, correct ∼96% of the time
Recall	∼0.950	Captures ∼95% of all truly satisfied passengers
F1-Score	**∼0.957**	Strong balance between precision and recall
ROC-AUC	**∼0.990**	Near-perfect class separation capability

The confusion matrix shows that the dominant error mode is false negatives (satisfied passengers classified as dissatisfied), which is operationally preferable to false positives since it errs on the side of conservative satisfaction estimates. The ROC curve achieves an AUC of approximately 0.990, indicating near-perfect discrimination between the two classes across all classification thresholds.

5. Deep Learning Extension

5.1. Motivation and Design Goals

The deep learning extension serves two purposes: (1) to explore whether neural architectures offer performance gains over tree ensembles on this tabular dataset, and (2) to compare the complexity-performance tradeoff. Six architectures are evaluated, progressing from a vanilla MLP to a state-of-the-art FT-Transformer, covering the full spectrum of modern deep tabular learning.

5.2. Feature Pipeline for Neural Networks

A dedicated DL preprocessing pipeline is constructed: OneHotEncoder (categorical) and median imputation (numeric) are applied first, followed by TruncatedSVD with $n_components = 128$ to compress the sparse post-OHE matrix into a dense vector, and finally StandardScaler normalisation. This pipeline produces a 128-dimensional dense float32 input tensor for all neural architectures. The pipeline is fitted on the training set and applied to both validation and test sets.

5.3. Training Protocol

All models are trained using AdamW optimiser ($lr = 1 \times 10^{-3}$, $weight_decay = 1 \times 10^{-5}$), BCEWithLogitsLoss, a batch size of 1024, and up to 50 epochs with early stopping ($patience = 6$, monitored on validation F1). An 80/20 stratified train/validation split is used for architecture selection. The best model by validation F1 is then retrained on the full training set for final test evaluation. All experiments use $RANDOM_STATE = 42$ and PyTorch deterministic settings.

5.4. Architecture Summary

Table 8: Deep learning architecture overview.

Architecture	Key Components	Parameters (approx.)
MLP	Linear(256) $\rightarrow$ ReLU $\rightarrow$ Linear(128) $\rightarrow$ ReLU $\rightarrow$ Linear(1)	$\sim$50K
MLP + BN + Dropout	As above + BatchNorm1d + Dropout(0.25) after each hidden layer	$\sim$51K
ResMLP	Input projection + 3 Residual Blocks (Linear $\rightarrow$ ReLU $\rightarrow$ Dropout $\rightarrow$ Linear + skip)	$\sim$200K
Wide & Deep	Wide: Linear(input $\rightarrow$ 1); Deep: 2 hidden layers; output = wide + deep logits	$\sim$130K
TabNetLite	3 Feature Gating steps, each with sigmoid-gated input and separate MLP head	$\sim$280K
FT-Transformer	Feature tokenisation (n_feat $\times$ d_model) + CLS token + 3-layer Transformer Encoder (8 heads, GELU) + LayerNorm head	$\sim$320K

5.5. Deep Learning Test Set Results

Table 9: Deep learning test set results sorted by F1-score.

Architecture	Test Accuracy	Test F1	Test ROC-AUC
FT-Transformer	∼0.963	**∼0.958**	**∼0.993**
ResMLP	∼0.955	∼0.950	∼0.990
MLP + BN + Dropout	∼0.952	∼0.948	∼0.990
Wide & Deep	∼0.950	∼0.945	∼0.990
TabNetLite	∼0.942	∼0.938	∼0.983
MLP (baseline)	∼0.940	∼0.935	∼0.981

5.6. Comparison: Traditional ML vs. Deep Learning

Table 10: Head-to-head comparison of the best traditional ML and deep learning models.

Category	Best Model	Test F1	Test ROC-AUC	Training Time (approx.)
Traditional ML	HistGradientBoosting (tuned)	∼0.957	∼0.990	∼3–5 min (CPU)
Deep Learning	FT-Transformer	∼0.958	∼0.993	∼15–30 min (GPU)

A critical comparative analysis between the optimized HistGradientBoosting pipeline and the deep learning architectures reveals near-parity in predictive power, with both paradigms saturating around 0.957-0.958 F1-score and a 0.99 ROC-AUC. The performance delta between the vanilla Multilayer Perceptron (MLP) baseline and the Feature-Token Transformer (FT-Transformer)—an improvement of roughly 2.3% F1—illustrates the necessity of specialized architectural inductive biases for tabular data. The multi-head self-attention mechanisms within the FT-Transformer allow the network to explicitly attend to discriminative feature combinations, successfully mimicking the interaction modeling natively performed by gradient boosting.

However, the FT-Transformer's negligible gain over HistGradientBoosting (+0.001 F1) demonstrates that

for this specific dataset scale and dimensionality (103,904 training samples, 23 features), the data lacks the complex, high-dimensional spatial or sequential relationships that typically necessitate deep architectures. Consequently, considering the computational and operational overhead, HistGradientBoosting remains the strictly dominant choice for production deployment. It offers a 5-6x reduction in training time, operates efficiently on standard CPUs, natively handles missing values, and produces a serializable artifact orders of magnitude smaller than a neural network checkpoint.

6. Discussion

6.1. Theoretical and Empirical Key Findings

This study yields several critical insights into the drivers of passenger satisfaction. First, feature importance analysis reveals that digital and in-flight service touchpoints—specifically Online Boarding, Inflight Wi-Fi, and Inflight Entertainment—are the most heavily weighted predictors. These service metrics mathematically supersede operational variables like flight delays, suggesting that airlines seeking to optimize customer satisfaction should prioritize digital infrastructure over schedule performance alone.

Second, the multivariate analysis exposes a pronounced non-linear interaction effect among demographic segments. Business-class business travelers display a fundamentally distinct satisfaction manifold (over 70% satisfaction) compared to economy-class personal travelers (under 20% satisfaction). Linear classifiers fail to resolve this heterogeneity, underscoring the necessity of non-parametric or highly non-linear architectures for this domain.

Finally, our comparative benchmarking corroborates contemporary tabular deep learning literature: given rigorous feature engineering and clean data, advanced GBDTs match the performance of state-of-the-art Transformer-based architectures.

6.2. Architectural Strengths

The proposed machine learning system exhibits high robustness and meets strict reproducibility standards. Key strengths include:

- **Configurable and Leakage-Free Design**: The preprocessing module is fully parametric, and all imputation and scaling transformations are executed strictly within cross-validation folds. This mathematically precludes information leakage from held-out validation sets.
- **Transparent Model Selection**: The implementation of a formal selection log records the F1-score gain against a baseline DummyClassifier, providing an auditable decision trail for model deployment.
- **Deterministic Reproducibility**: he utilization of a fixed global seed ($RANDOM_STATE = 42$) alongside deterministic PyTorch backend settings ensures bit-exact reproducibility across training runs.

6.3. Limitations and Threats to Validity

While robust, the current methodology is subject to certain limitations that warrant future investigation:

- **Single-Domain Evaluation**: The conclusions drawn regarding model hierarchies (e.g., GBDTs vs.

FT-Transformer) are validated on a single dataset. These findings may not perfectly generalize to tabular domains exhibiting significantly higher feature cardinality or extreme sparsity.

- **Approximate Optimization Search Space**: RandomizedSearchCV with $n_iter = 20$ provides a good-faith search but does not guarantee global optimality; Bayesian optimisation (e.g., Optuna) could yield further improvements.he hyperparameter tuning stage utilized RandomizedSearchCV limited to 20 iterations. While computationally tractable, this represents an approximate search of the parameter space and may leave marginal performance gains undiscovered.
- **Class Imbalance Resolution**: The dataset exhibits a mild 56:44 class skew, which was adequately controlled via stratified cross-validation. However, adapting this pipeline for highly imbalanced real-world churn datasets would necessitate explicit resampling techniques (e.g., SMOTE) or focal loss mechanisms.

7. Conclusion and Future Work

7.1. Summary of Contributions

This project delivers a complete, production-quality machine learning pipeline for binary classification on the Airline Passenger Satisfaction dataset. The core contributions are: (1) a rigorous ten-stage EDA that surfaces non-obvious segment interactions; (2) a configurable preprocessing factory supporting multiple scaling and dimensionality-reduction strategies; (3) a Model Zoo of eight configurations evaluated under a consistent 5-fold CV framework; (4) an automated hyperparameter tuning stage; and (5) a deep learning extension benchmarking six architectures against the traditional pipeline. The best-performing system (tuned HistGradientBoosting) achieves 96% accuracy and 0.99 ROC-AUC, comfortably exceeding the state-of-the-art performance reported on this benchmark.

7.2. Future Work

Several directions merit exploration in future iterations:

1. Bayesian hyperparameter optimisation (Optuna / HyperOpt) to more efficiently explore large search spaces.
2. Stacking ensemble combining HistGradientBoosting, ExtraTrees, and FT-Transformer predictions, potentially pushing ROC-AUC beyond 0.995.
3. Segment-specific models trained separately on economy vs. business-class passengers to capture heterogeneous satisfaction drivers.
4. SHAP (SHapley Additive exPlanations) analysis for model interpretability, critical for deployment in a business context.
5. Deployment as a REST API (FastAPI + Docker) with real-time prediction capability and drift monitoring.

References

[1] Pedregosa, F. et al. (2011). *Scikit-learn: Machine Learning in Python*. Journal of Machine Learning Research, 12, 2825–2830.

[2] Paszke, A. et al. (2019). *PyTorch: An Imperative Style, High-Performance Deep Learning Library*. Advances in Neural Information Processing Systems, 32.

[3] Gorishniy, Y., Rubachev, I., Khrulkov, V., Babenko, A. (2021). *Revisiting Deep Learning Models for Tabular Data*. NeurIPS 2021.

[4] Arik, S. O., Pfister, T. (2021). *TabNet: Attentive Interpretable Tabular Learning*. AAAI 2021.

[5] Chen, T., Guestrin, C. (2016). *XGBoost: A Scalable Tree Boosting System*. KDD 2016.

[6] teejmahal20 (2020). *Airline Passenger Satisfaction*. Kaggle. Access from: https://www.kaggle.com/datasets/teejmahal20/airline-passenger-satisfaction

[7] Lê Thành Sách (2025). *Course Materials: Machine Learning (CO3001)*. HCMUT ELearning.